

MOHAN NALAMOTHU

Mobile: [REDACTED]

E-Mail: [REDACTED]

LinkedIn- <https://www.linkedin.com/in/mohan-nalamothu-62168224a/>



Career Objective

Seeking a challenging role in a dynamic environment where I can apply and enhance my technical skills in data engineering, cloud computing, and software development. Looking to contribute to an organization that fosters professional growth and innovation.

Professional Summary

Experienced Senior Data Engineer 11 years of total IT experience with proficient in Snowflake, AWS with expertise in data modeling, ETL development, and data analytics. Skilled in Snowflake streams, tasks, time travel, Python, Kafka, AWS Glue, AWS Kinesis, Airflow, Lambda, Step functions, Azure Blob, Azure Data Factory (ADF), AWS SQS, Azure Queue Services and more. Proven track record of delivering successful projects across various domains including Insurance, Banking, Sales, Telecom.

Roles and Responsibilities

- Build and maintain Enterprise Data Platforms to support analytics tools.
- Design and implement data migration process from existing Data Warehouse to Cloud Data Warehouses.
- Implement and maintain modern data ETL/ELT pipelines to pull data from heterogeneous source systems.
- Built customized ETL framework using Go.
- Design data quality and governance frameworks to ensure data accuracy and consistency.
- Create security policies and role-based access controls to share and access the data.
- Design data pipeline job monitoring frameworks with event driven architecture.
- Experienced in handling different file formats like Text, CSV, Avro, JSON, Parquet etc.
- Hands on experience in various enterprise data platform phases like Data Ingestion, Data Processing and Data Analytics.
- In-depth knowledge of Snowflake Database, Multi-cluster architecture and credit Usage.
- Handled incremental data using Snowflake streams & tasks.

- Have good experience on Snowflake Failsafe, Time travel and Zero-Copy cloning.
- In-Depth understanding & experience of clustering to improve compute performance.
- Experienced in building Snowpipe for continuous data ingestion from the S3 bucket.
- Experience in working with the AWS lambda using Python (Boto3) snowflake connector to push large sized data to the snowflake from the AWS S3 buckets.
- Skilled in optimization and performance tuning of enterprise data platforms.
- Experience in leveraging AWS Kinesis Data Streams for robust data pipeline architectures, facilitating continuous data capture and immediate processing for analytics and decision-making applications.
- Used AWS Glue, ADF for efficient server less data transformations, ensuring optimized data workflows, and integrating seamlessly with Snowflake analytic services for comprehensive data solutions.
- Have good knowledge of Azure Data Lake, Databricks architecture and functionalities.
- Experienced in schema design and dimensional data modelling.
- Knowledge of Data Warehousing concepts and Dimensional Data. Fact and Dimensions tables, Slowly Changing Dimensions (SCD's).
- Strong usage knowledge of CI, CD tools like Azure DevOps and Jenkins.
- Experience in Creating Database and Database objects like Tables, Schema, Stored Procedures, Views, Functions etc.
- Design and develop interactive visualizations using Tableau, PowerBI and AWS QuickSight.
- Experienced with Agile development process and practices.

CERTIFICATIONS

Snowflake Certified: SnowPro Core Certification

EDUCATION

Master of Computer Application (MCA) TN, India, 2013

Technical Skills:

Snowflake Services	Snow SQL, Externalstages, SnowPIPE, Snowpark, Integrations, Streams, Tasks, Sharing, TimeTravel, Zero-CopyCloning, Clustering, Multi-Cluster Warehousing.
AWS Services	EC2, S3, Glue, Lambda, AWS Glue
Databases	Snowflake, Teradata, Oracle, SQL Server, MySQL
Scripting	Python
IDE and Developer Tools	Jupyter Notebook, Anaconda, VS Code and GitHub CoPilot

Version Control&CI/CD	GIT
Monitoring Tools	Airflow, Autosys, Control-M

Projects Handled

Project #1: GBI Finance

Company: Cognizant Technology Solutions

Sr. Associate (Sep 2021 – Feb 2025)

Client: Apple

Role: Sr. Associate

Environment: Snowflake, Kafka, Keystone, SQL, Python, Autosys, GitHub, Airflow, Oracle, Big Query, Tableau

Description:

GBI Finance is a project where we load the client details to snowflake cloud from multiple third-party teams for business reporting.

- Developed ETL pipelines for data extraction, transformation, and loading (ETL) from Kafka, SAP, and S3 to Snowflake.
- Implemented query optimization techniques to improve performance.
- Built data validation scripts using Python.
- Designed email notification mechanisms for automated failure alerts.
- Created dashboards and reports using Tableau.

Project #2: Terminus Migration

Client: Apple

Role: Sr. Associate

Environment: Snowflake, Teradata, Kafka, Keystone, SQL, Python, Autosys, GitHub, Airflow

Description:

Migration of ETL's from Teradata to Snowflake which sourced from multiple sources like HANA, ORACLE etc. This project was part of bigger program for Client's cloud transition.

- Led the successful migration of a large-scale Teradata data warehouse to Snowflake, enabling cloud scalability and reducing operational costs significantly.
- Designed and executed a phased migration strategy, transferring numerous tables and views from Teradata to Snowflake with zero downtime.

- Converted complex Teradata SQL scripts to Snowflake-compatible syntax, optimizing multiple stored procedures for improved performance.
- Implemented Snowflake's Search Optimization Service, enhancing query response times significantly for critical analytical workloads.
- Applied clustering techniques in Snowflake to optimize data retrieval, cutting scan times noticeably on high-volume tables.
- Automated DDL and data migration processes using a suitable tool/language, reducing manual effort considerably while ensuring data integrity.
- Validated and reconciled massive volumes of historical data post-migration, ensuring complete accuracy between Teradata and Snowflake environments.
- Integrated Snowflake with an ETL tool, streamlining downstream data pipelines and enabling real-time analytics for multiple users.
- Collaborated with cross-functional teams to retire legacy Teradata infrastructure, completing the cutover ahead of schedule.
- Optimized Snowflake resource allocation and compute costs, achieving a notable reduction in monthly cloud expenses compared to Teradata.
- Developed and executed comprehensive validation scripts to compare Teradata and Snowflake data, achieving full row-level consistency across critical tables.
- Created automated reconciliation dashboards in a tool Tableau to monitor data quality post-migration, reducing validation time and ensuring stakeholder confidence.
- Migrated and optimized Teradata table reports to Snowflake, leveraging dynamic SQL and Time Travel to enhance reporting flexibility and reduce refresh times.
- Redesigned legacy Teradata reporting tables into Snowflake's schema, enabling seamless integration with a BI tool and improving report generation speed.

Project #3: Nordea Credit

Company: Capgemini solutions India Private limited

Sr. Associate (Sep 2018 – Sep 2021)

Client: Nordea

Role: Sr Consultant

Environment: Snowflake, Teradata, Python, Unix, Oracle, Tivoli, GitHub

Responsibilities:

- Involved in development in Core and SBSL level as per the client's requirements.
- Analyzing the data movement from Source to stage followed by core and SBSL.
- Working on Procedures and Views to load the data from stage to core and SBSL based on the logics given in mapping document.
- Moving large amounts of data into and out of other AWS data stores to snowflake, such as Amazon Simple Storage Service (Amazon S3) using Airflow custom objects.
- Creation of Unit test cases and implementation of unit test cases in HP QC.

- Created Changing Dimension SCD type-2 & Implemented Stage delta and Target delta.
- After development checking the data flow with the help of unit testing.
- Expertise in code deployment in higher environments via GIT tool.
- Have worked on production Defects for every release raised in core and SBSL layers.
- Prepared Monthly release reports to get the UAT sign off from business for my Team.
- Involved in Project Transition to the AM team by preparing the handover plan with respective documentation.

Project #4: RDW AUTO

Company: Capgemini solutions India Private limited

Client: ANZ

Role: Sr Consultant

Environment: Teradata SQL Assistant, Putty, WinSCP, CTRL-M

Project Summary:

RDW – it is a data analytics and reporting platform and supports regulatory/operational reporting across ANZ businesses in Asia with implemented countries including Taiwan, Philippines, Hong Kong, India and Japan.

- RDW uses connect Direct to receive files from Core banking systems like FCL, MUREX, MIDANZ and people soft and loads Data into staging area.
- Data flow of ETL in RDW -- once the files are received from business team into the interface area, it undergoes the formatting process (when required) and then goes through transformation and finally loaded into PDM.
- RDW run on window platform to receive files and then connects to Teradata installed on UNIX box to load the data.
- RDW follow FSLDM model for transformation.
- In RDW we are using an in-built tool called DW Automation on Windows AU server which in turn connects to Teradata and ETL. This is used for job scheduling, sequencing and also to create the job dependency.
- Teradata utilities like BTEQ, Fast load, Fast exports are also runs on the same DW Automation server. Final extracts exported from Teradata are generated on the window server and sends to down streams using connect Direct.
- Connect Direct jobs are scheduled through control-m in cyclic which monitor for extracts and sends when generated.

Project #5: Sales and Analysis

Company: Aarya technical consulting private limited

Sr. Associate (Feb 2014 – Sep 2018)

Client: Vodafone

Role: Associate Engineer

Environment: Teradata SQL Assistant, Putty, WinSCP, CTRL-M

Project Summary:

- Architect/Design – Prepare the Design Document by using data provided by Architect team or using solution document. If we have any doubts in solution document, we can take clarification on them.
- Coding – Modifications in existing data extraction process for extracting data from source for a faster and reliable extraction, Teradata Fast load script to load data into landing area. BTEQ script to populate data into target tables through which Facts and Dimensions tables are built for generating reports for clients.
- Involved in performance tuning and Error handling. Creating various tables like volatile and global temporary tables and creating indexes for query optimization.
- Unit Testing – Complete testing considering the scenarios that could rise in future as well as the scenarios that are mentioned in the requirements.

Project #5: New England Health Plan

Company: Aarya technical consulting private limited

Client: Cardinal Health

Role: Teradata Developer

Environment: Teradata 15, UNIX, Teradata SQL and Informatica

Project Summary:

Cardinal Health is one of the leading health Insurance provider organizations in USA. The scope of the New England Health Plan (NEHP) track is to develop foundational-build requirements necessary to support Large and Small Group accounts with NEHP products migrating from the New England States (CT, ME, NH) to WGS. The NEHP work is only intended to take on product specific build in order to be able to administer group enrollments, claims processing, servicing, etc.

- We will get BSD, based on this document we prepare FL, ML, BTEQ scripts based on requirement.
- Preparing the Technical Design document (TTD) based up on the given BSD document.
- Workload Management (WLM) creation and Database objects (DBOBJECTS), Self-authorization form preparation.
- For created WLM jobs Migrating BTEQ Script or Informatica workflow from DEV to SIT and SIT to IR and IR to EDWARD.
- Running the created job sets through The WLM tool.
- Preparing the CAT LOGS and DBQL LOGS for scripts and informatics mappings.
- Creating DB objects like tables, views based up on the requirement.
- Working with various Indexes such as PI, SI etc.
- Identifying long running Queries, Scripts and tuning them.

- Creating volatile, Derived, Global temporary tables based on requirement.
- Did unit testing and prepare unit test cases and we capture the results.
- Extensively worked with UNIX commands like grep, copy, zip, unzip, sed, cut, archive, remove, compare etc.
- Supporting while the job sets are running in production, after completing the production loads Handover the entire project to LO team.

Key Competencies

- ✓ Data Engineering & Big Data Technologies
- ✓ Cloud Data Migration (Snowflake)
- ✓ ETL & Data Warehousing
- ✓ Python, SQL & Performance Tuning
- ✓ SDLC & Agile Methodologies
- ✓ DevOps & Automation
- ✓ Stakeholder & Team Collaboration
- ✓ Production Support & Incident Management